

# Contents

## Units

- Layers [Convolution]

- Layers [Convolution] [Receptive field]

- Layers [Dilated Convolution, Deformable Convolution]

- Layers [Upsampling, Learnable Upsampling]

- Layers [Batch Norm, Dropout]

- Layers [Group Convolutions and its variants]

## Blocks

- VGG

- Inception

- ResNet\*

- MobileNet\*

## Architectures


- VGG, Inception, ResNet\*, MobileNet\*

## Neural Architecture Search (NAS)

## Summary

# Big Picture

## Net




# Big Picture


Net



Blocks



# Big Picture



# Contents

## Units

### Layers [Convolution]

Layers [Convolution] [Receptive field]

Layers [Dilated Convolution, Deformable Convolution]

Layers [Upsampling, Learnable Upsampling]

Layers [Batch Norm, Dropout]

Layers [Group Convolutions and its variants]

## Blocks

VGG

Inception

ResNet\*

MobileNet\*


## Architectures

VGG, Inception, ResNet\*, MobileNet\*

## Neural Architecture Search (NAS)


## Summary

# Layers [Convolution]




# Layers [Convolution]

- kernel size  $F$  ?
- padding size  $P$  ?
- stride  $S$  ?




# Layers [Convolution]

- kernel size  $F = 4 \times 4$
- padding size  $P = 2$
- stride  $S = 1$






# Layers [Convolution]



- kernel size,  $F = 3 \times 3$
- padding size,  $P = 1$
- stride,  $S = 2$

# Layers [Convolution]



- kernel size,  $F = 3 \times 3$
- padding size,  $P = 1$
- stride,  $S = 2$

# Contents

## Units

- Layers [Convolution]

- Layers [Convolution] [Receptive field]

- Layers [Dilated Convolution, Deformable Convolution]

- Layers [Upsampling, Learnable Upsampling]

- Layers [Batch Norm, Dropout]

- Layers [Group Convolutions and its variants]

## Blocks

- VGG

- Inception

- ResNet\*

- MobileNet\*


## Architectures

- VGG, Inception, ResNet\*, MobileNet\*

## Neural Architecture Search (NAS)

## Summary

# Layers [Convolution] [Receptive field]



- kernel size,  $F = 3 \times 3$
- padding size,  $P = 1$
- stride,  $S = 2$



The **receptive field** is defined as the region in the input space that a particular CNN's feature is looking at (i.e. be affected by)

A receptive field of a feature can be described by

- its center location
- its size


# Layers [Convolution] [Receptive field]

- kernel size  $F = 3 \times 3$ ,
- padding size  $P = 1$ ,
- stride  $S = 2$



# Layers [Convolution] [Receptive field]


- kernel size  $F = 3 \times 3$ ,
- padding size  $P = 1$ ,
- stride  $S = 2$



# Layers [Convolution] [Receptive field]


Layer 0:

$$n_0 = 5; r_0 = 1; j_0 = 1; \\ start_0 = 0.5$$




Conv1:

$$k_1 = 3; p_1 = 1; s_1 = 2$$




Layer 1:

$$n_1 = 3; r_1 = 3; j_1 = 2; \\ start_1 = 0.5$$




Layer 1:

$$n_1 = 3; r_1 = 3; j_1 = 2; \\ start_1 = 0.5$$




Conv2:

$$k_2 = 3; p_2 = 1; s_2 = 2$$



Layer 2:

$$n_2 = 2; r_2 = 7; j_2 = 4; \\ start_2 = 0.5$$



$$n_{out} = \left\lfloor \frac{n_{in} + 2p - k}{s} \right\rfloor + 1$$

$$j_{out} = j_{in} * s$$

$$r_{out} = r_{in} + (k - 1) * j_{in}$$

$$start_{out} = start_{in} + \left( \frac{k - 1}{2} - p \right) * j_{in}$$

# Contents

## Units

- Layers [Convolution]

- Layers [Convolution] [Receptive field]

- Layers [Dilated Convolution, Deformable Convolution]

- Layers [Upsampling, Learnable Upsampling]

- Layers [Batch Norm, Dropout]

- Layers [Group Convolutions and its variants]

## Blocks

- VGG

- Inception

- ResNet\*

- MobileNet\*

## Architectures

- VGG, Inception, ResNet\*, MobileNet\*

## Neural Architecture Search (NAS)

## Summary




# Layers [Dilated Convolution]

(used for *Semantic Segmentation*, *Object Recognition*)

Also known as “atrous convolutions”.

Dilated convolutions “inflate” the kernel by inserting spaces between the kernel elements. The dilation “rate” is controlled by an additional hyperparameter **d**.

- input size  $I = 7 \times 7$
- kernel size  $F = 3 \times 3$ ,
- padding size  $P = 0$ ,
- stride  $S = 1$
- dilation rate **d** = 2



# Layers [Dilated Convolution]

(used for *Semantic Segmentation*, *Object Recognition*)

- input size  $I = 7 \times 7$
- kernel size  $F = 3 \times 3$ ,
- padding size  $P = 0$ ,
- stride  $S = 1$
- dilation rate  $d = 2$

For 'normal' convolution we have:


$$W2 = (W1 - F + 2P) / S + 1,$$

$$H2 = (H1 - F + 2P) / S + 1$$

$$D2 = K$$

$$F' = (F - 1) \cdot (d - 1)$$


How this will change ?



# Layers [Dilated Convolution]


(used for *Semantic Segmentation*, *Object Recognition*)

- Detection of fine-details by processing inputs in higher resolutions.
- Broader view of the input to capture more contextual information.
- Faster run-time with less parameters



# Layers [Deformable Convolution]

(used for *Semantic Segmentation*, *Object Recognition*)




$$\mathcal{R} = \{(-1, -1), (-1, 0), \dots, (0, 1), (1, 1)\}$$

$$y(\mathbf{p}_0) = \sum_{\mathbf{p}_n \in \mathcal{R}} w(\mathbf{p}_n) \cdot \mathbf{x}(\mathbf{p}_0 + \mathbf{p}_n)$$

# Layers [Deformable Convolution]

(used for *Semantic Segmentation*, *Object Recognition*)

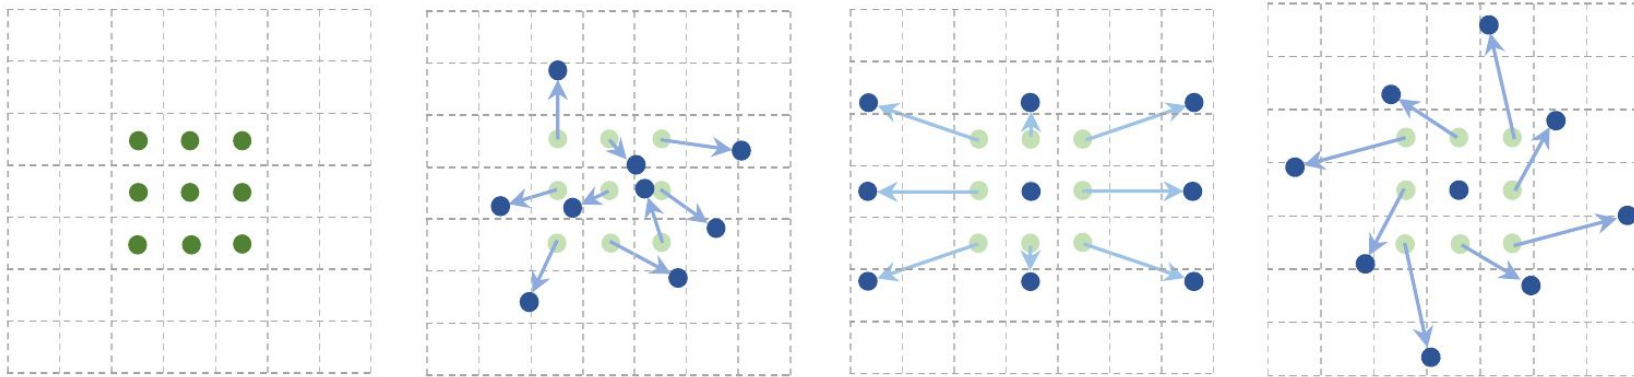


$$\mathcal{R} = \{(-1, -1), (-1, 0), \dots, (0, 1), (1, 1)\}$$

$$y(\mathbf{p}_0) = \sum_{\mathbf{p}_n \in \mathcal{R}} w(\mathbf{p}_n) \cdot \mathbf{x}(\mathbf{p}_0 + \mathbf{p}_n) \quad \longrightarrow \quad y(\mathbf{p}_0) = \sum_{\mathbf{p}_n \in \mathcal{R}} w(\mathbf{p}_n) \cdot \mathbf{x}(\mathbf{p}_0 + \mathbf{p}_n + \Delta \mathbf{p}_n)$$

# Layers [Deformable Convolution]

(used for *Semantic Segmentation*, *Object Recognition*)




$$\mathcal{R} = \{(-1, -1), (-1, 0), \dots, (0, 1), (1, 1)\}$$


$$y(\mathbf{p}_0) = \sum_{\mathbf{p}_n \in \mathcal{R}} w(\mathbf{p}_n) \cdot \mathbf{x}(\mathbf{p}_0 + \mathbf{p}_n) \quad \longrightarrow \quad y(\mathbf{p}_0) = \sum_{\mathbf{p}_n \in \mathcal{R}} w(\mathbf{p}_n) \cdot \mathbf{x}(\mathbf{p}_0 + \mathbf{p}_n + \Delta \mathbf{p}_n)$$

# Layers [Deformable Convolution]

(used for *Semantic Segmentation*, *Object Recognition*)



# Layers [Deformable Convolution]




(a) standard convolution

(b) deformable convolution



# Layers [Deformable Convolution]



# Contents

## Units

- Layers [Convolution]

- Layers [Convolution] [Receptive field]

- Layers [Dilated Convolution, Deformable Convolution]

- Layers [Upsampling, Learnable Upsampling]

- Layers [Batch Norm, Dropout]

- Layers [Group Convolutions and its variants]

## Blocks

- VGG

- Inception

- ResNet\*

- MobileNet\*

## Architectures

- VGG, Inception, ResNet\*, MobileNet\*

## Neural Architecture Search (NAS)

## Summary

# Layers [Upsampling]

- Nearest Neighbor

1	2
3	4



1	1	2	2
1	1	2	2
3	3	4	4
3	3	4	4

Input: [2 x 2]

Output: [2 x 2]

- Bilinear

1	2
3	4




1.00	1.25	1.75	2.00
1.50	1.75	2.00	2.50
2.50	2.75	3.25	3.50
3.00	3.25	3.75	4.00

Input: [2 x 2]


Output: [2 x 2]

# Layers [Learnable Upsampling: Transpose Convolution]



# Layers [Learnable Upsampling: Transpose Convolution]

[ Heavily used for *Semantic Segmentation*, *GANs*, *Autoencoders* ]



# Contents

## Units

Layers [Convolution]

Layers [Convolution] [Receptive field]

Layers [Dilated Convolution, Deformable Convolution]

Layers [Upsampling, Learnable Upsampling]

Layers [Dropout, Batch Norm]

Layers [Group Convolutions and its variants]

## Blocks

VGG

Inception

ResNet\*

MobileNet\*

## Architectures

VGG, Inception, ResNet\*, MobileNet\*


## Neural Architecture Search (NAS)

## AutoML


## Summary

# Layers [Dropout]

- In-network ensembling
- Reduce overfitting



(a) Standard Neural Net




(b) After applying dropout.



# Layers [Batch Normalization]

BN: data-driven normalizing each layer, for each batch

- Greatly accelerate training
- Less sensitive to initialization
- Improve regularization




- $\mu$ : mean of  $x$  in mini-batch
- $\sigma$ : std of  $x$  in mini-batch
- $\gamma$ : scale
- $\beta$ : shift

- $\mu, \sigma$ : functions of  $x$ , analogous to responses
- $\gamma, \beta$ : parameters to be learned, analogous to weights



# Layers [Batch Normalization]



# Contents

## Units

Layers [Convolution]

Layers [Convolution] [Receptive field]

Layers [Dilated Convolution, Deformable Convolution]

Layers [Upsampling, Learnable Upsampling]

Layers [Batch Norm, Dropout]

Layers [Group Convolutions and its variants]

## Blocks

VGG

Inception

ResNet\*

MobileNet\*

## Architectures

VGG, Inception, ResNet\*, MobileNet\*



## Neural Architecture Search (NAS)

## Summary

# Layers [Group Convolution]

Normal  
Convolution:




conv 5x5



# Layers [Group Convolution]

Normal  
Convolution:

conv 3x3





# Layers [Group Convolution]

Normal  
Convolution:



conv 1x1

A special name:  
point-wise  
convolutions




# Layers [Group Convolution]

gconv 3x3 (g=2)




# Layers [Group Convolution]

gconv 3x3 (g=3)




# Layers [Group Convolution]

gconv 3x3 (g=3)




special name -  
**depth-wise**  
**convolutions** or  
**channel-wise**  
**convolutions**





# Layers [Group Convolution]

gconv 1x1 (g=3)



# Contents

## Units

- Layers [Convolution]

- Layers [Convolution] [Receptive field]

- Layers [Dilated Convolution, Deformable Convolution]

- Layers [Upsampling, Learnable Upsampling]

- Layers [Batch Norm, Dropout]

- Layers [Group Convolutions and its variants]

## Blocks

- VGG

- Inception

- ResNet\*

- MobileNet\*


## Architectures

- VGG, Inception, ResNet\*, MobileNet\*

## Neural Architecture Search (NAS)

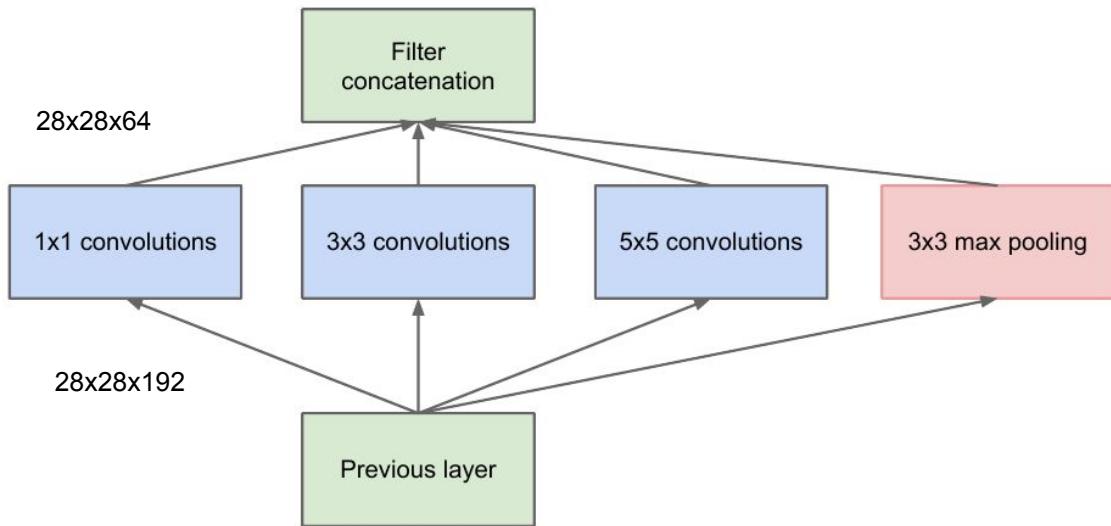
## Summary

# Blocks [VGG]



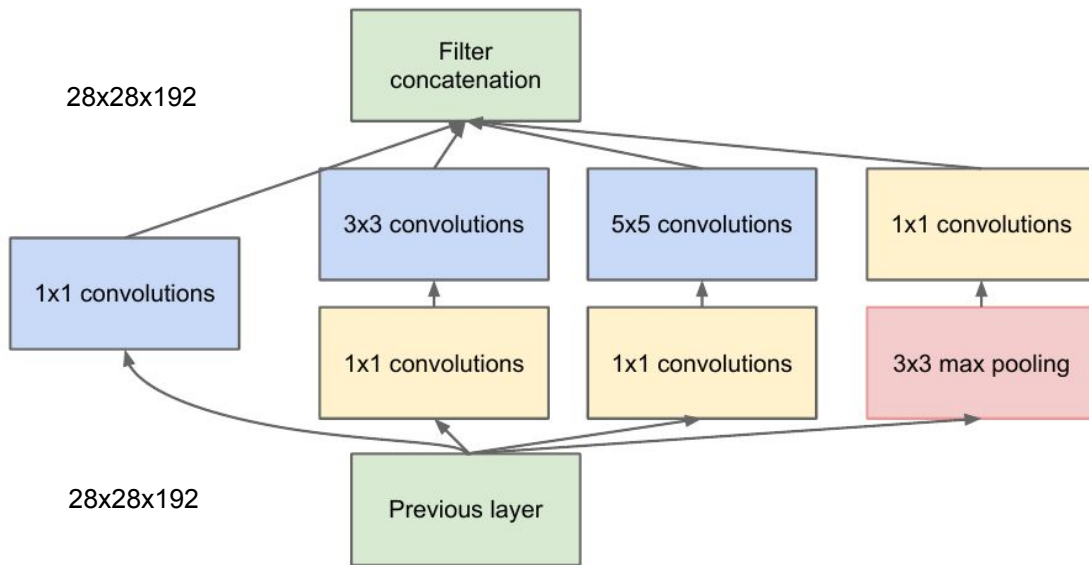
# Blocks [GoogleNet / Inception]

- to find out optimal local **sparse structure** and to repeat it spatially
- to split operations for cross-channel correlations and at spatial correlations into a series of independently operations.
- split-transform-merge strategy



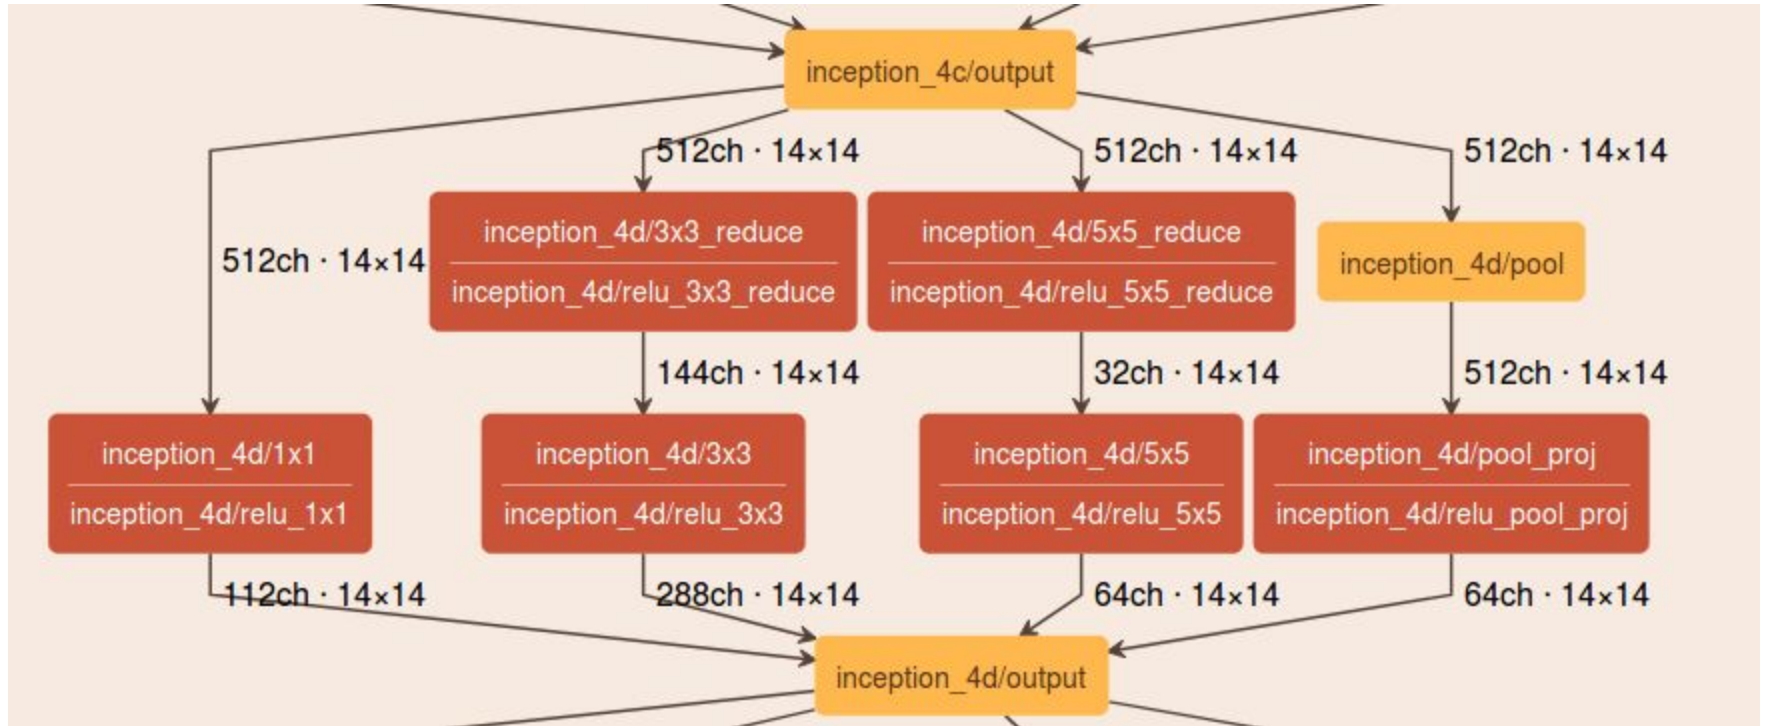
# Blocks [GoogleNet / Inception]

- to find out optimal local sparse structure and to repeat it spatially
- to split operations for cross-channel correlations and at spatial correlations into a series of independently operations.
- split-transform-merge strategy



## Bottleneck


# Blocks [GoogleNet / Inception]




# Blocks [ResNet]

$$G(x) = x + F(x)$$

In the basic design,  $F(x)$  contains two  $3 \times 3$  convolution layers along with a batch normalization and/or a rectified linear unit activation function.




# Blocks [ResNet, bottleneck]



For deeper networks  
(ResNet-50+), use “bottleneck”  
layer to improve efficiency  
(similar to Inception)





# Blocks [ResNet, bottleneck]




conv 1x1

conv 3x3


conv 1x1




# Blocks [ResNet, bottleneck]




=



# Blocks [ResNeXt]





**Inception:**  
heterogeneous multi-branch



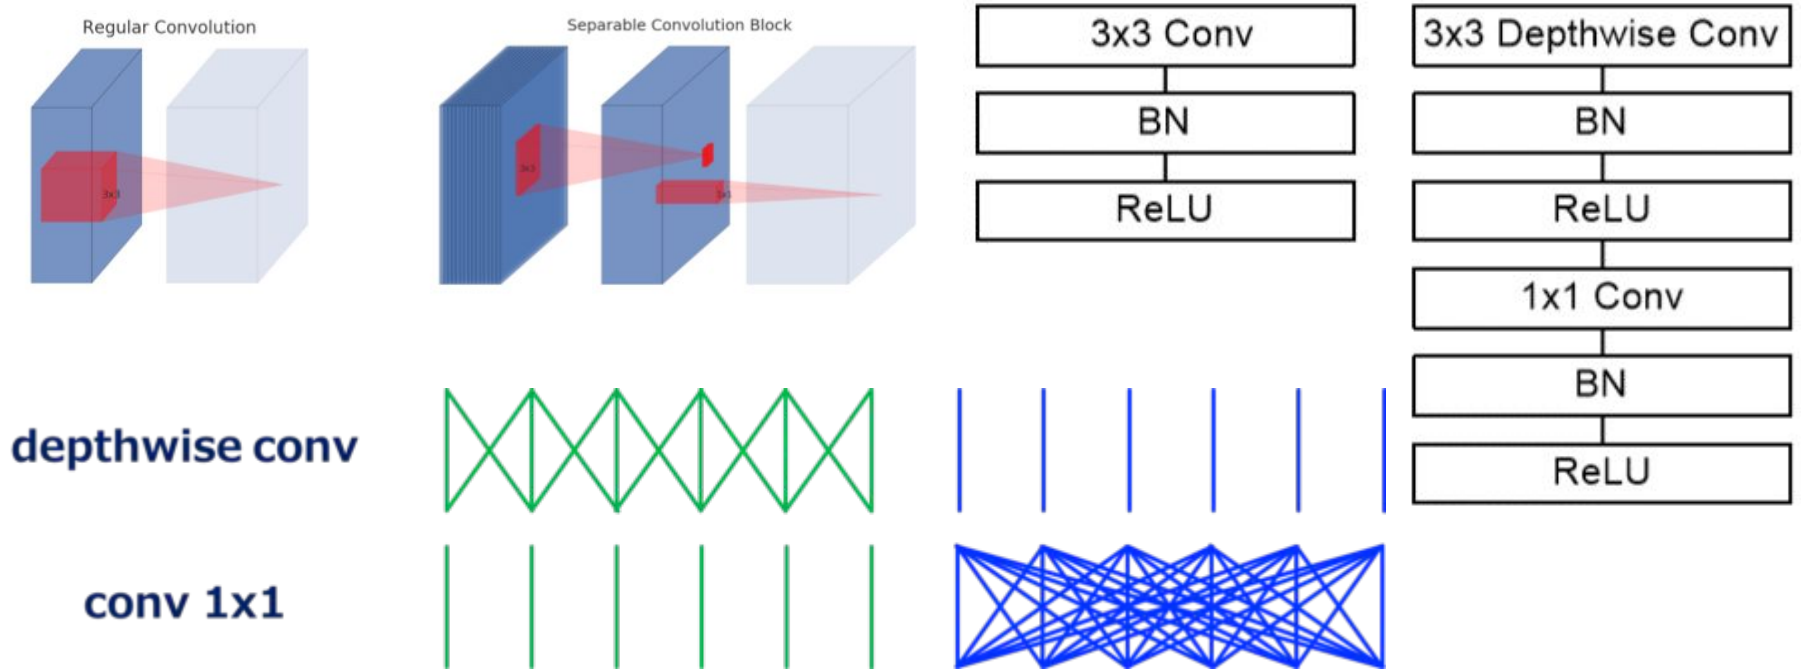
**ResNeXt:**  
uniform multi-branch

# Blocks [ResNeXt]



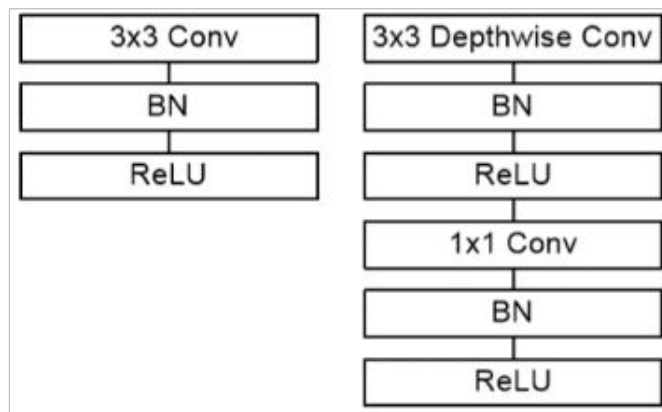
**ResNeXt:**  
uniform multi-branch

# Blocks [MobileNet V1]



# MobileNet v1

- $M$  – number of input channels
- $N$  – the number of output channels
- $D_K * D_K$  the kernel size
- $D_F * D_F$  the feature map size



Standard convolution computation cost

$$D_K \cdot D_K \cdot M \cdot N \cdot D_F \cdot D_F$$

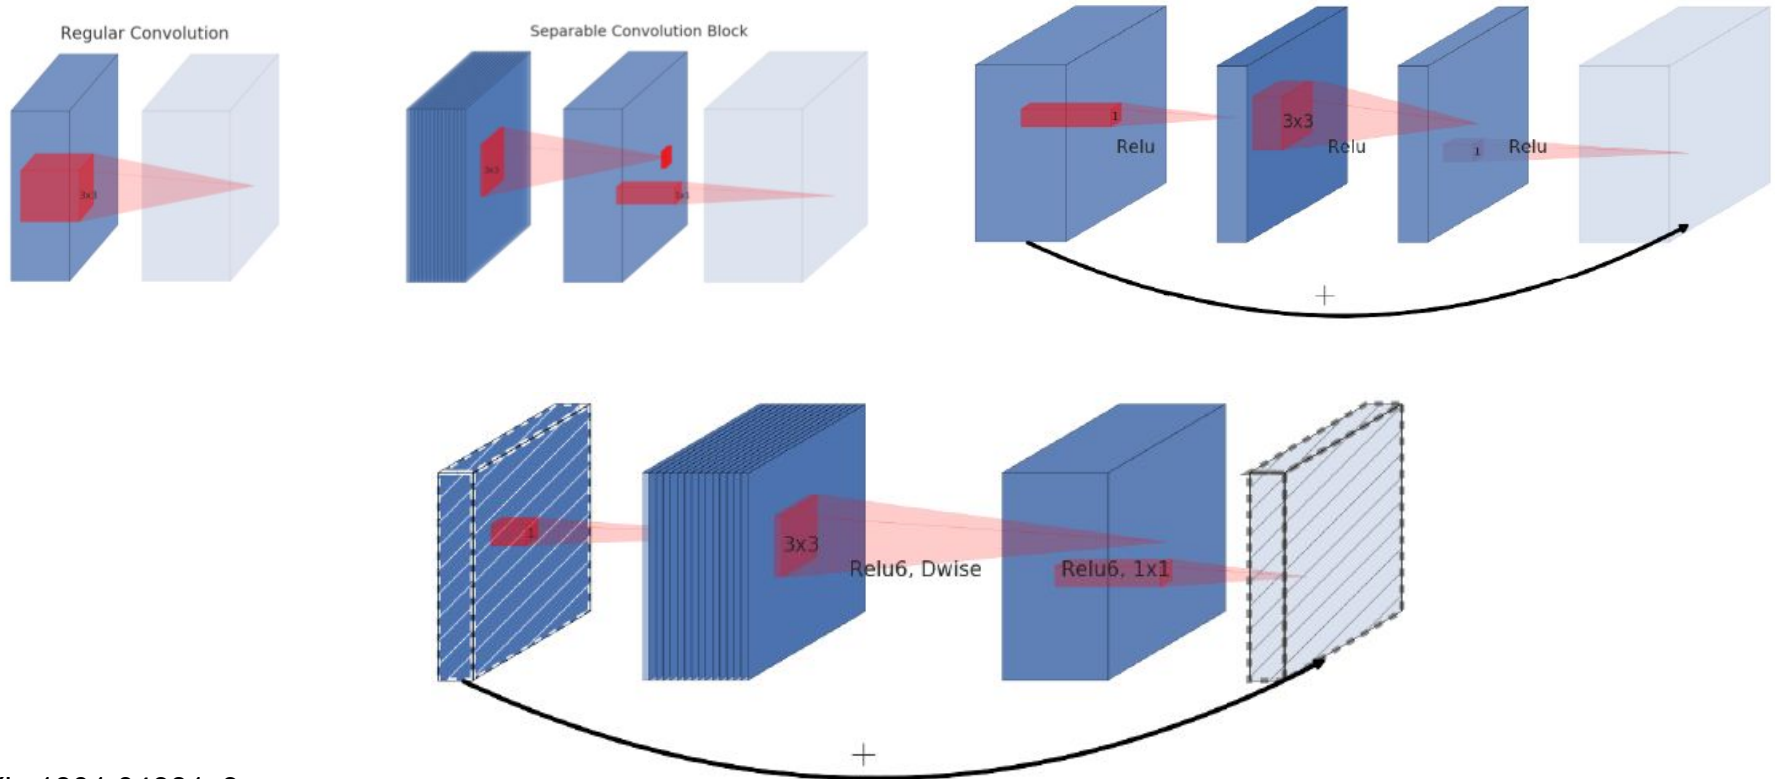
Depthwise separable convolution computational cost

$$D_K \cdot D_K \cdot M \cdot D_F \cdot D_F + M \cdot N \cdot D_F \cdot D_F$$

Reduction in computation:

$$\frac{D_K \cdot D_K \cdot M \cdot D_F \cdot D_F + M \cdot N \cdot D_F \cdot D_F}{D_K \cdot D_K \cdot M \cdot N \cdot D_F \cdot D_F} = \frac{1}{N} + \frac{1}{D_K^2}$$


# Blocks [MobileNet V2]



arXiv:1801.04381v3

<https://medium.com/@yu4u/why-mobilenet-and-its-variants-e-g-shufflenet-are-fast-1c7048b9618d>

# Blocks [MobileNet V2]



arXiv:1801.04381v3

<https://medium.com/@yu4u/why-mobilenet-and-its-variants-e-g-shufflenet-are-fast-1c7048b9618d>



# Contents

## Units

- Layers [Convolution]

- Layers [Convolution] [Receptive field]

- Layers [Dilated Convolution, Deformable Convolution]

- Layers [Upsampling, Learnable Upsampling]

- Layers [Batch Norm, Dropout]

- Layers [Group Convolutions and its variants]

## Blocks

- VGG

- Inception

- ResNet\*

- MobileNet\*


## Architectures

- VGG, Inception, ResNet\*, MobileNet\*

## Neural Architecture Search (NAS)


## Summary

# Net [VGG16]

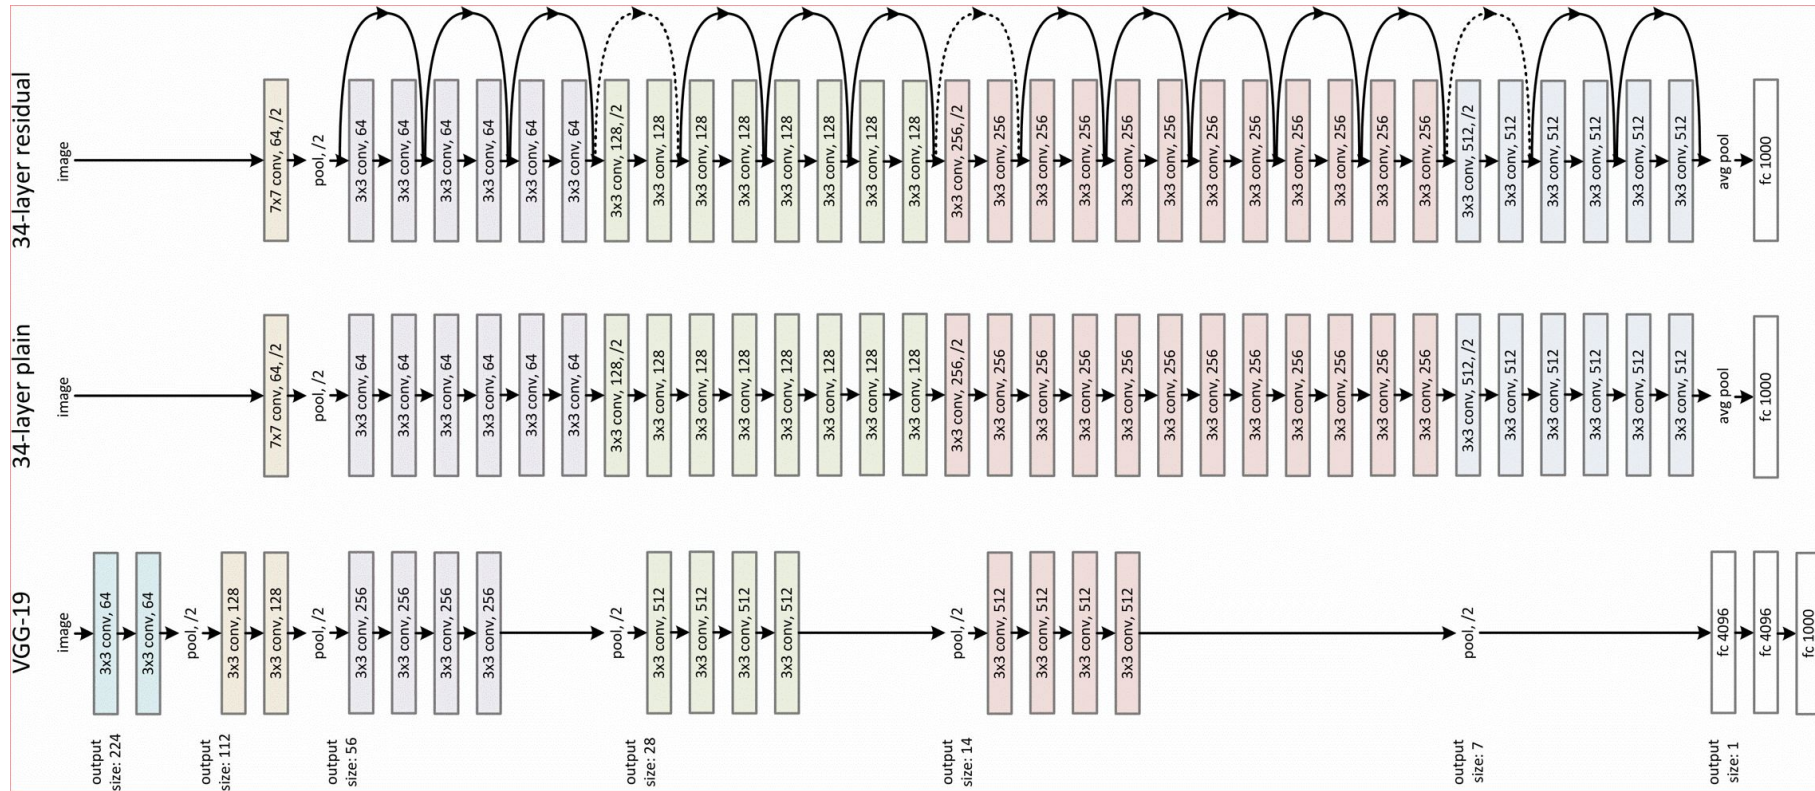


- 3 3x3 Conv as the module
- Stack the same module
- Same computation for each module  
(1/2 spatial size => 2x filters)

# Net [GoogLeNet]



# Net [ResNet & ResNetX]

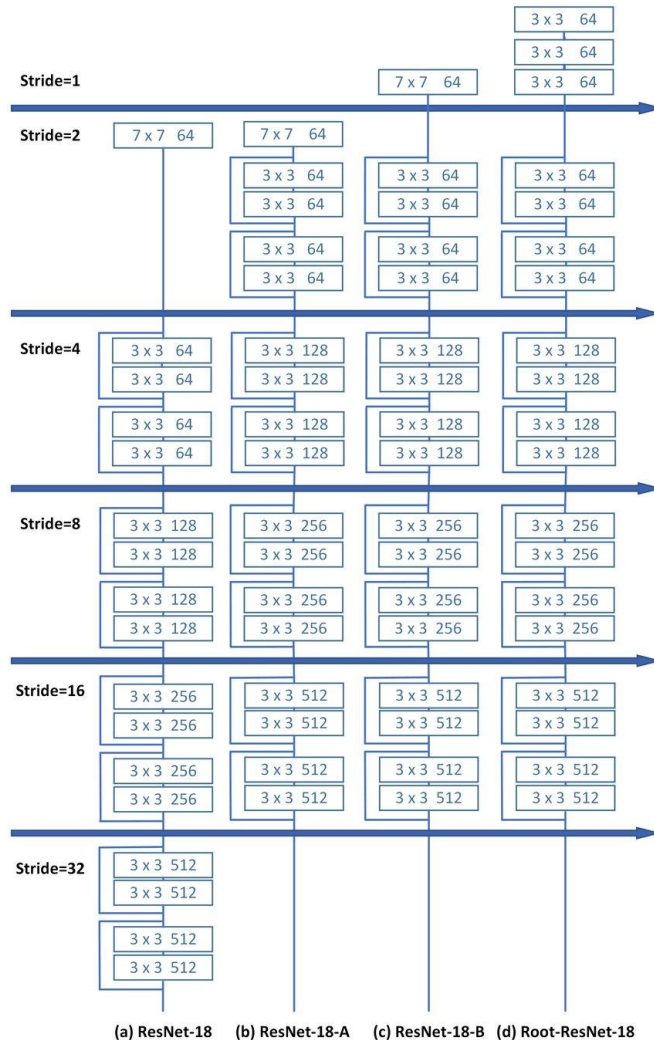


# Net [ResNet]

layer name	output size	18-layer	34-layer	50-layer	101-layer	152-layer
conv1	$112 \times 112$	$7 \times 7$ , 64, stride 2				
conv2_x	$56 \times 56$	$3 \times 3$ max pool, stride 2				
		$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$
conv3_x	$28 \times 28$	$\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 8$
conv4_x	$14 \times 14$	$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 23$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 36$
conv5_x	$7 \times 7$	$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$
	$1 \times 1$	average pool, 1000-d fc, softmax				
FLOPs		$1.8 \times 10^9$	$3.6 \times 10^9$	$3.8 \times 10^9$	$7.6 \times 10^9$	$11.3 \times 10^9$

# Net [ResNet]

- (a) ResNet18: original structure.
  - (b) ResNet-18-A: removing the first maxpooling layer.
  - (c) ResNet-18-B: changing the stride size in the first conv layer from 2 to 1.
  - (d) Root-ResNet-18: replacing the  $7 \times 7$  conv layer with three stacked  $3 \times 3$  conv layers in ResNet18-B.
- The corresponding mAPs on PASCAL 2007 test (training on "07+12" from scratch) are 73.1%, 75.3%, 77.6% and 78.5%, respectively.





# Net [ResNetX]

Table 1. **(Left)** ResNet-50. **(Right)** ResNeXt-50 with a  $32\times 4d$  template (using the reformulation in Fig. 3(c)). Inside the brackets are the shape of a residual block, and outside the brackets is the number of stacked blocks on a stage. “ $C=32$ ” suggests grouped convolutions [24] with 32 groups. *The numbers of parameters and FLOPs are similar between these two models.*

stage	output	ResNet-50	ResNeXt-50 ( $32\times 4d$ )
conv1	$112\times 112$	$7\times 7$ , 64, stride 2	$7\times 7$ , 64, stride 2
conv2	$56\times 56$	$3\times 3$ max pool, stride 2	$3\times 3$ max pool, stride 2
		$\begin{bmatrix} 1\times 1, 64 \\ 3\times 3, 64 \\ 1\times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1\times 1, 128 \\ 3\times 3, 128, C=32 \\ 1\times 1, 256 \end{bmatrix} \times 3$
conv3	$28\times 28$	$\begin{bmatrix} 1\times 1, 128 \\ 3\times 3, 128 \\ 1\times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1\times 1, 256 \\ 3\times 3, 256, C=32 \\ 1\times 1, 512 \end{bmatrix} \times 4$
conv4	$14\times 14$	$\begin{bmatrix} 1\times 1, 256 \\ 3\times 3, 256 \\ 1\times 1, 1024 \end{bmatrix} \times 6$	$\begin{bmatrix} 1\times 1, 512 \\ 3\times 3, 512, C=32 \\ 1\times 1, 1024 \end{bmatrix} \times 6$
conv5	$7\times 7$	$\begin{bmatrix} 1\times 1, 512 \\ 3\times 3, 512 \\ 1\times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1\times 1, 1024 \\ 3\times 3, 1024, C=32 \\ 1\times 1, 2048 \end{bmatrix} \times 3$
	$1\times 1$	global average pool 1000-d fc, softmax	global average pool 1000-d fc, softmax
# params.		$25.5\times 10^6$	$25.0\times 10^6$
FLOPs		$4.1\times 10^9$	$4.2\times 10^9$

# Contents

## Units

- Layers [Convolution]

- Layers [Convolution] [Receptive field]

- Layers [Dilated Convolution, Deformable Convolution]

- Layers [Upsampling, Learnable Upsampling]

- Layers [Batch Norm, Dropout]

- Layers [Group Convolutions and its variants]

## Blocks

- VGG

- Inception

- ResNet\*

- MobileNet\*

## Architectures


- VGG, Inception, ResNet\*, MobileNet\*

Neural Architecture Search (NAS)

## Summary



# Neural Architecture Search (NAS)



# Neural Architecture Search (NAS)

- **searching for best blocks**

The is finding one particular building block which is then repeated many times to create the deep architecture.

- **starting from units**

Discover new connections of the whole net

# Neural Architecture Search (NAS)

start from units/simple blocks

- identity
- 1x7 then 7x1 convolution
- 3x3 average pooling
- 5x5 max pooling
- 1x1 convolution
- 3x3 depthwise-separable conv
- 7x7 depthwise-separable conv
- 1x3 then 3x1 convolution
- 3x3 dilated convolution
- 3x3 max pooling
- 7x7 max pooling
- 3x3 convolution
- 5x5 depthwise-separable conv

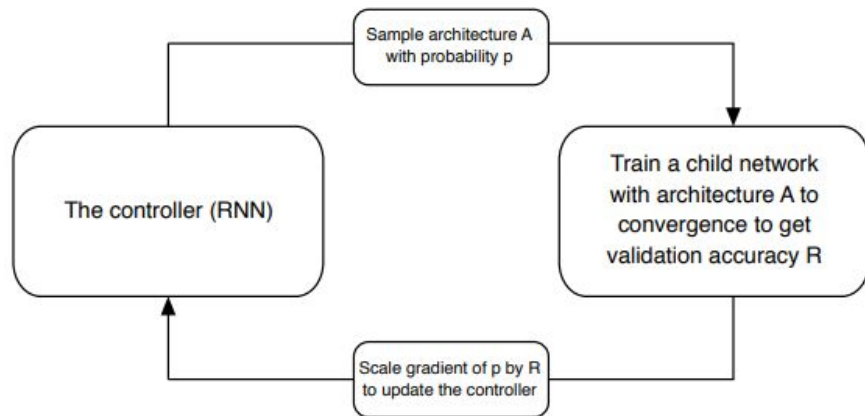
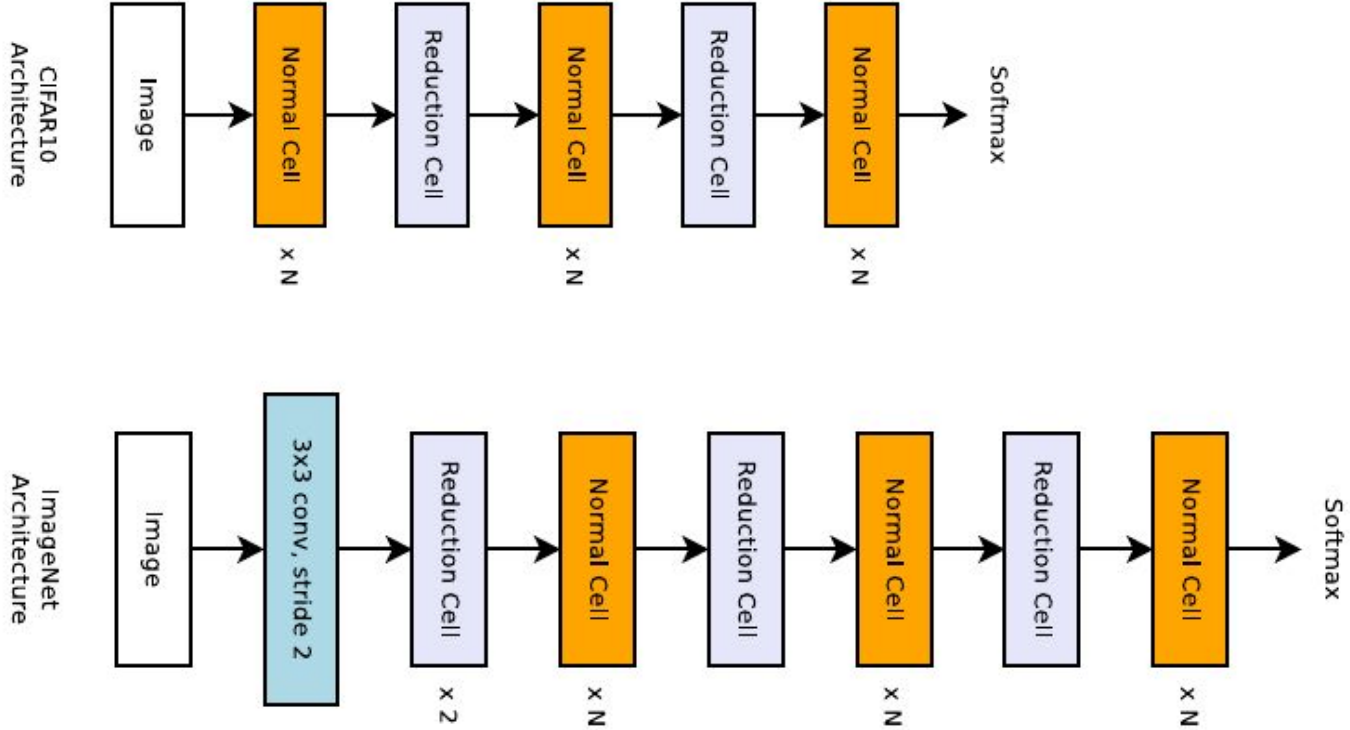
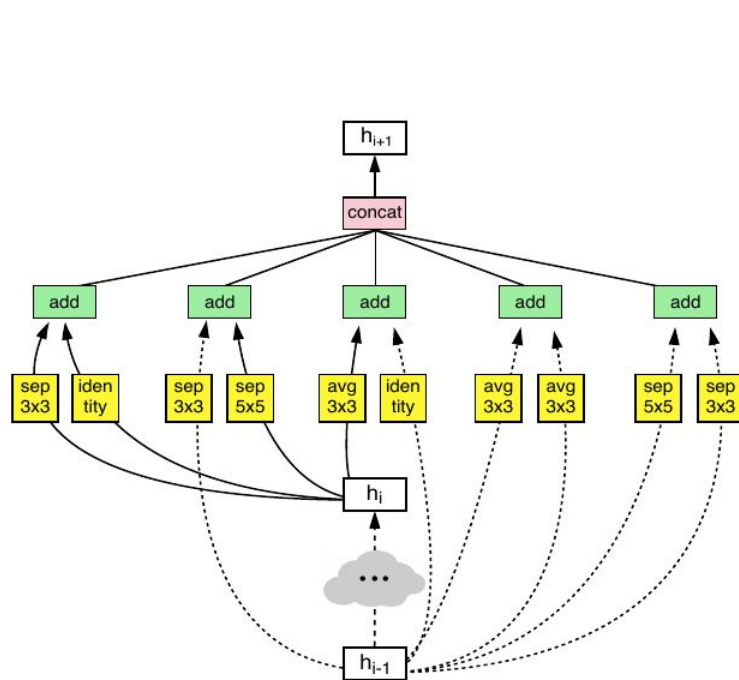


Figure 1. Overview of Neural Architecture Search [71]. A controller RNN predicts architecture  $A$  from a search space with probability  $p$ . A child network with architecture  $A$  is trained to convergence achieving accuracy  $R$ . Scale the gradients of  $p$  by  $R$  to update the RNN controller.

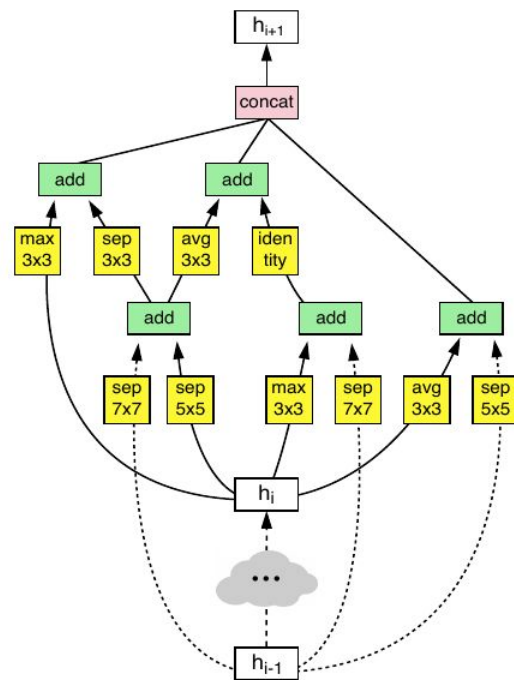
# Neural Architecture Search (NAS)



# Neural Architecture Search (NAS)



*Normal Cell*




*Reduction Cell*

# Neural Architecture Search (NAS)

Exploring Randomly Wired Neural  
Networks for Image Recognition  
[arXiv:1904.01569]

use three classical random graph  
models to generate randomly wired  
graphs for networks



# Contents

## Units

- Layers [Convolution]

- Layers [Convolution] [Receptive field]

- Layers [Dilated Convolution, Deformable Convolution]

- Layers [Upsampling, Learnable Upsampling]

- Layers [Batch Norm, Dropout]

- Layers [Group Convolutions and its variants]

## Blocks

- VGG

- Inception

- ResNet\*

- MobileNet\*

## Architectures


- VGG, Inception, ResNet\*, MobileNet\*

## Neural Architecture Search (NAS)

## Comparison

## Summary

# Net [Comparison]





# Contents

## Units

- Layers [Convolution]

- Layers [Convolution] [Receptive field]

- Layers [Dilated Convolution, Deformable Convolution]

- Layers [Upsampling, Learnable Upsampling]

- Layers [Batch Norm, Dropout]

- Layers [Group Convolutions and its variants]

## Blocks

- VGG

- Inception

- ResNet\*

- MobileNet\*

## Architectures

- VGG, Inception, ResNet\*, MobileNet\*

## Neural Architecture Search (NAS)

## Summary

# Summary

- When you see huge DN, don't be scare. Usually it can be decomposed.
- Automatic topology learning (NAS) is rising, but it is still important to understand basic blocks/units
- Classification problem is solved, but features matter